What is a Distributed File System?:

Why are there so many?


Education Records
Health Records
Auto Records
Auto Service
Career History
Other Deductions
DONATIONS
Interest Paid
Bank SAVINGS
Bank JOINT
Bank CHECKING
Credit Card
Retirement Account
Mortgage Statements
STATEMENTS, PAID BILLS, RECPT.

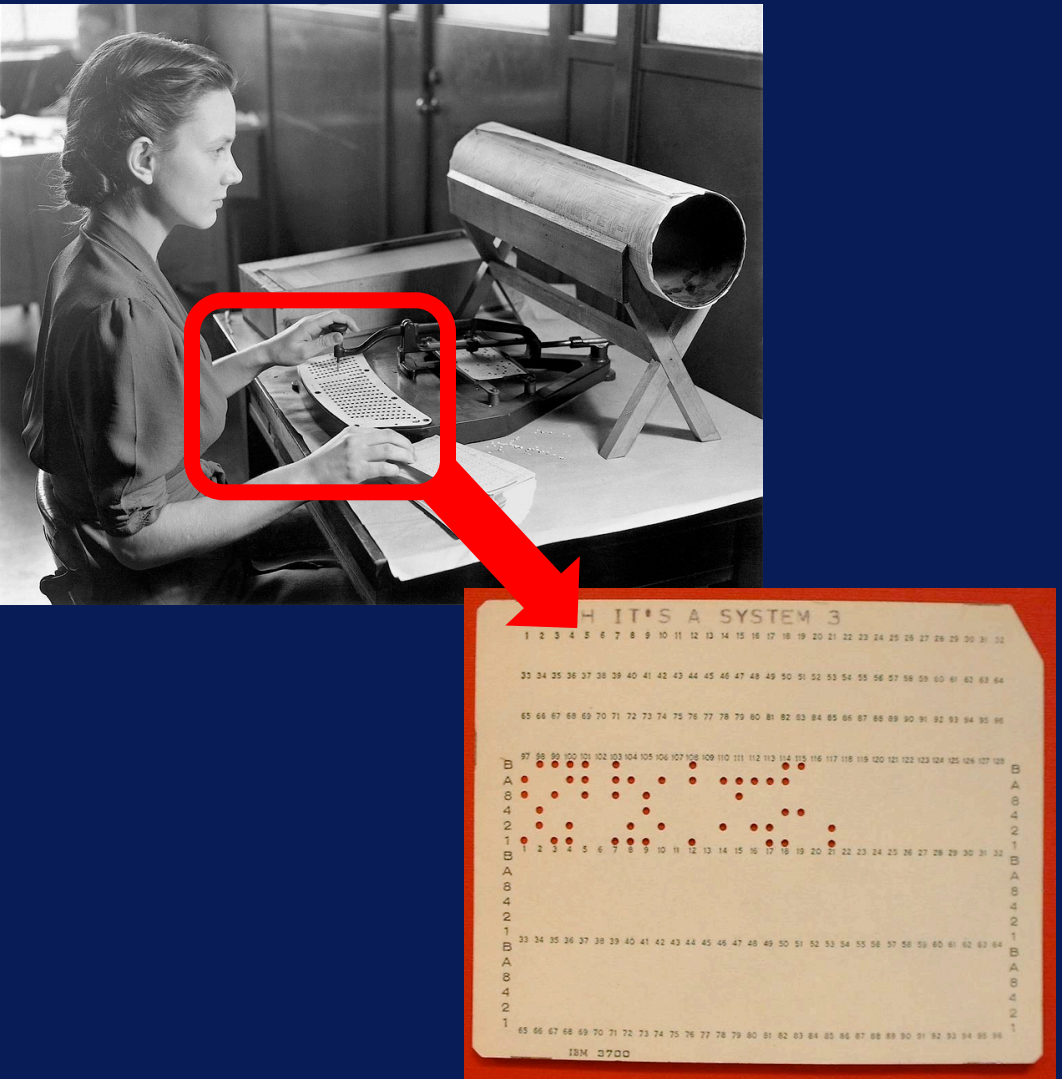
H IT'S A SYSTEM 3
IBM 3700

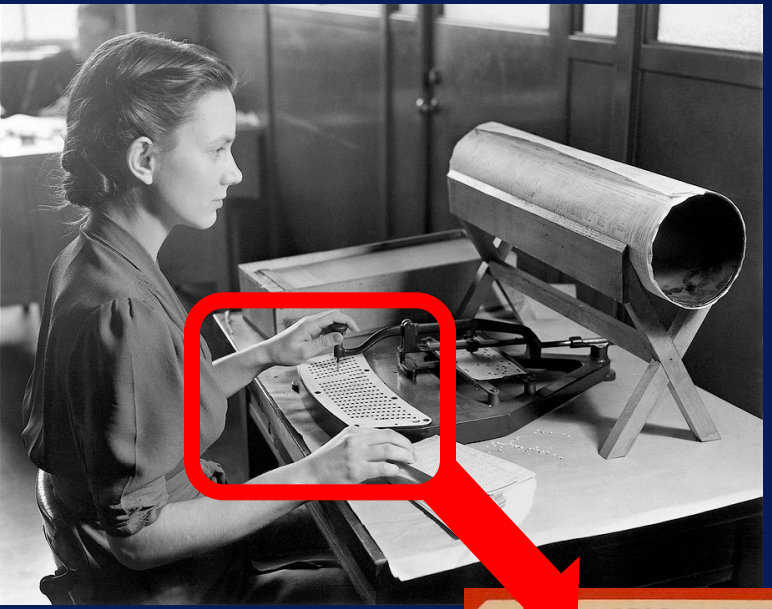

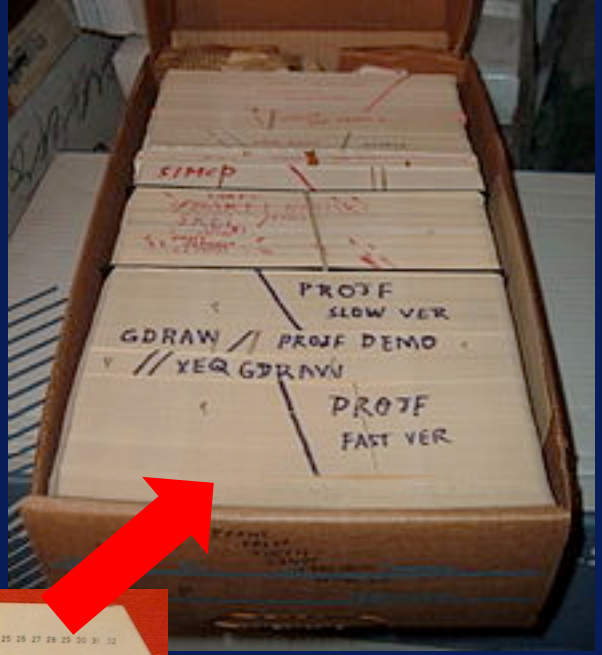
PROJF
SLOW VER
GDRAW // PROJF DEMO
// XEQ GDRAW
PROJF
FAST VER

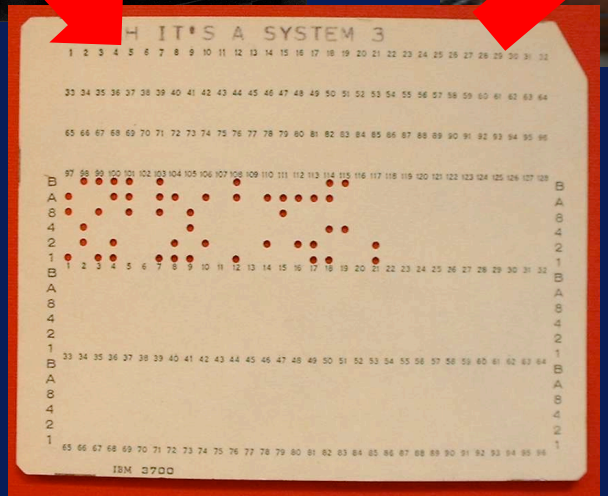
H IT'S A SYSTEM 3
IBM 3700

Long-term information storage

Long-term information storage

Access result of a process later

Long-term information storage

- Access result of a process later
- Store large amounts of information

Long-term information storage

- Access result of a process later
- Store large amounts of information
- Enable access of multiple processes

File

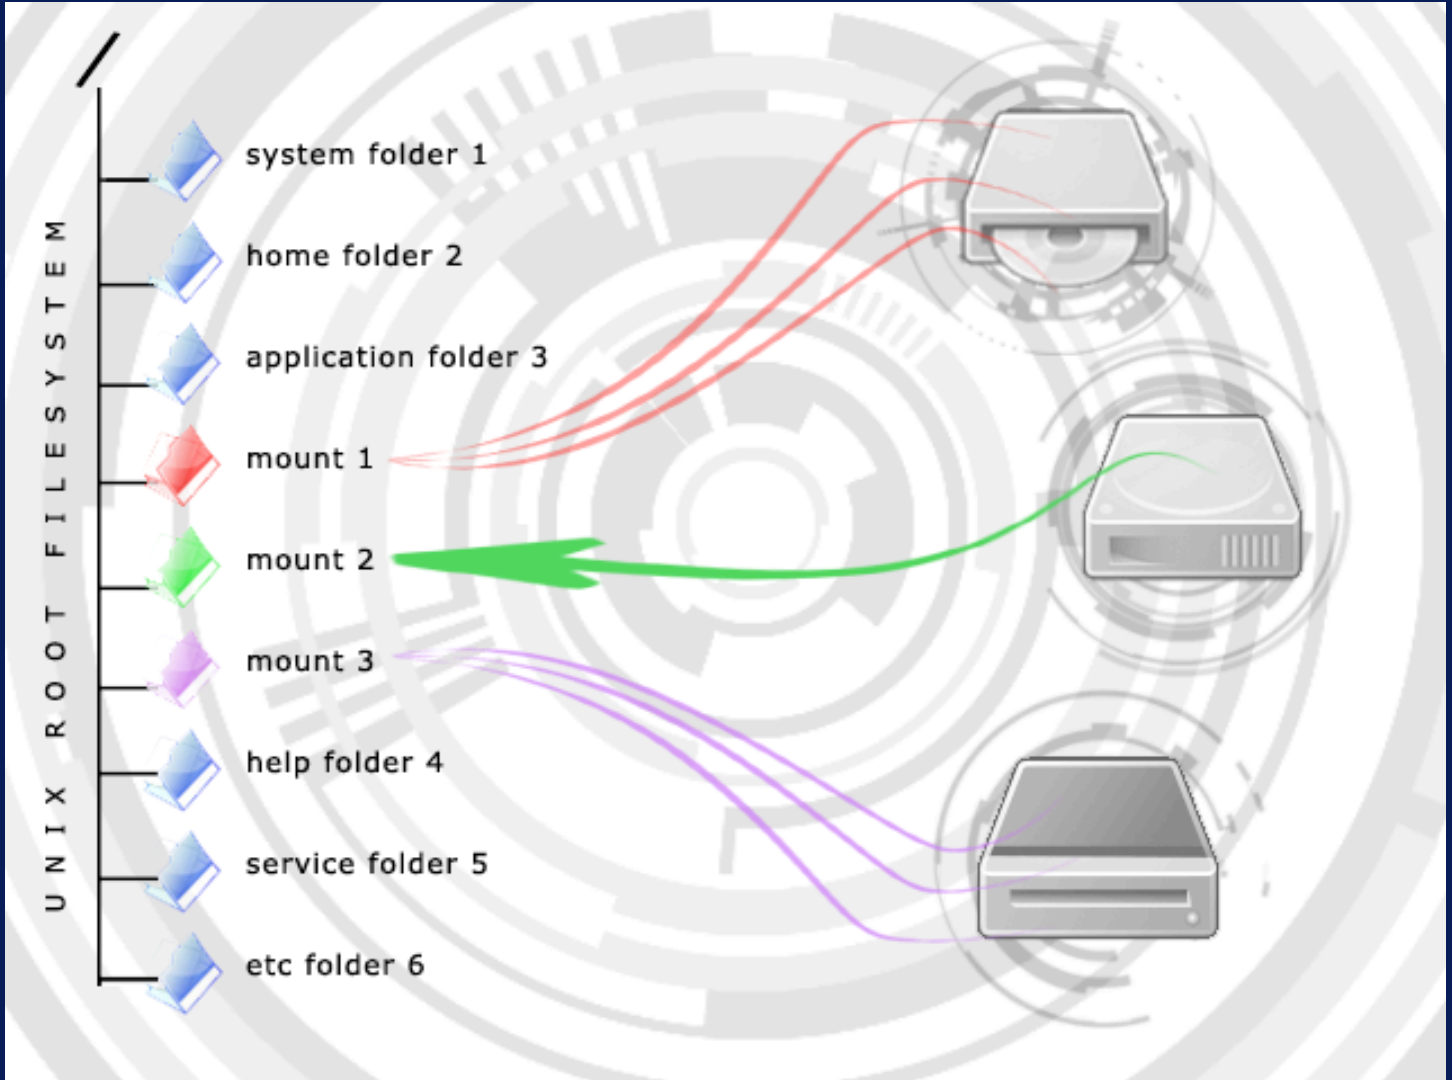
/
UNIX ROOT FILESYSTEM
system folder 1
home folder 2
application folder 3
mount 1
mount 2
mount 3
help folder 4
service folder 5
etc folder 6

Accessing files


64GB
256GB
1TB
5TB

? What if you have more data?

Buy a bigger disk?

OR

Copy data to an external hard drive?

WORK

PERSONAL

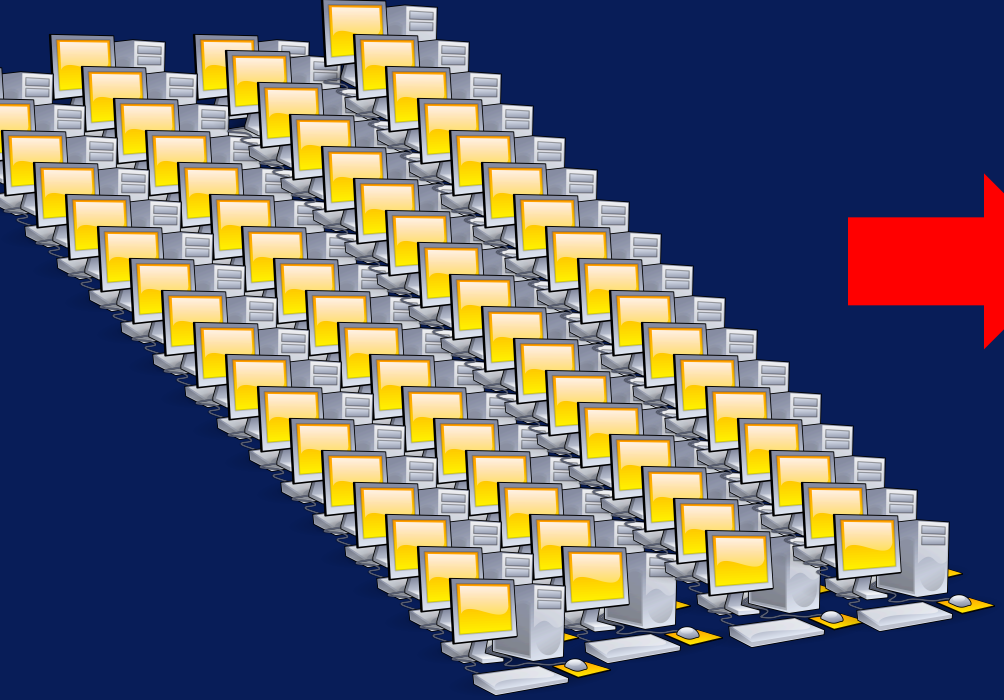

Rack

Rack

Distributed File System (DFS)

Data

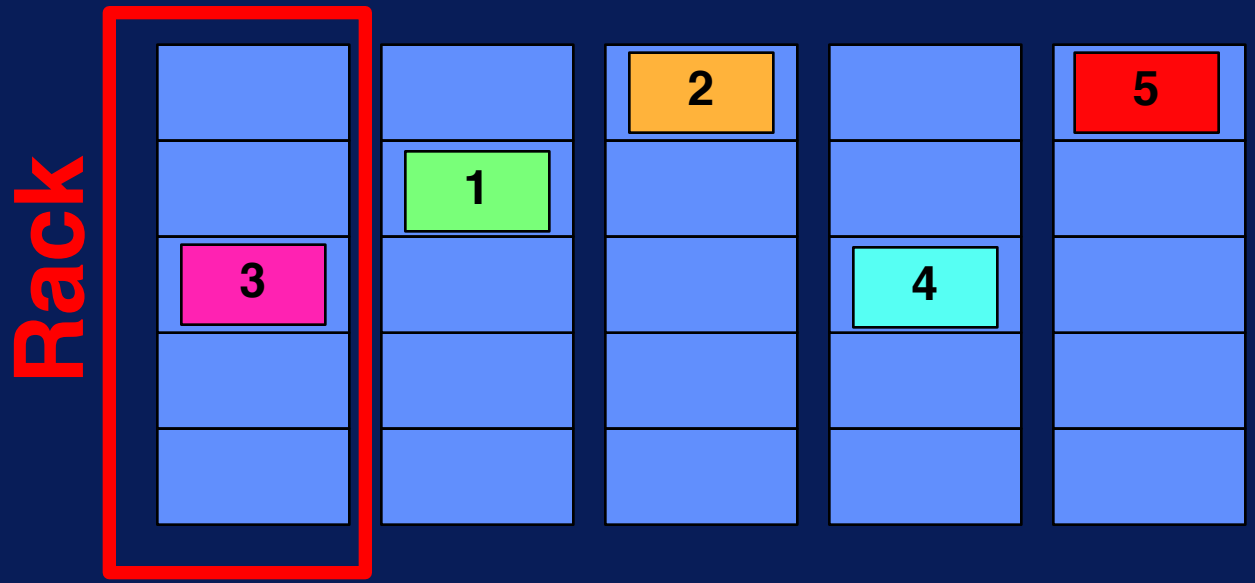

Data

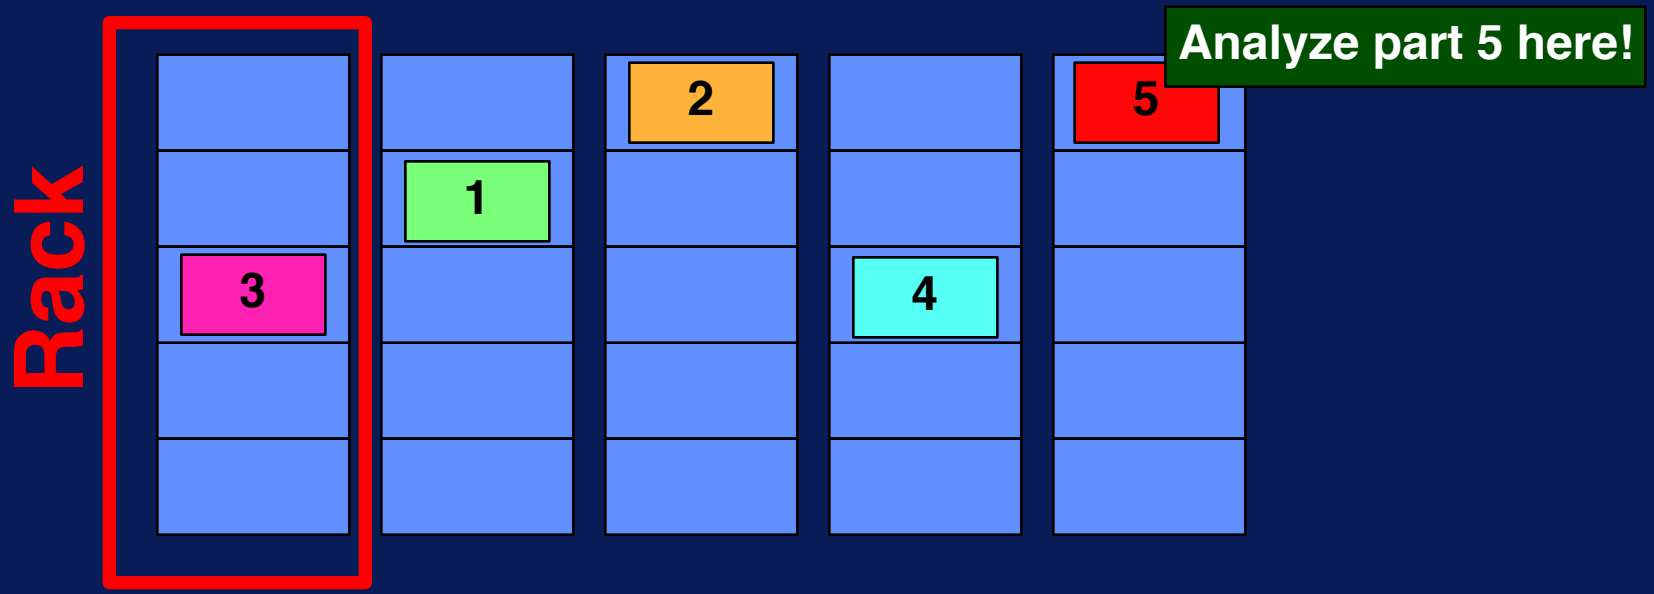

Data

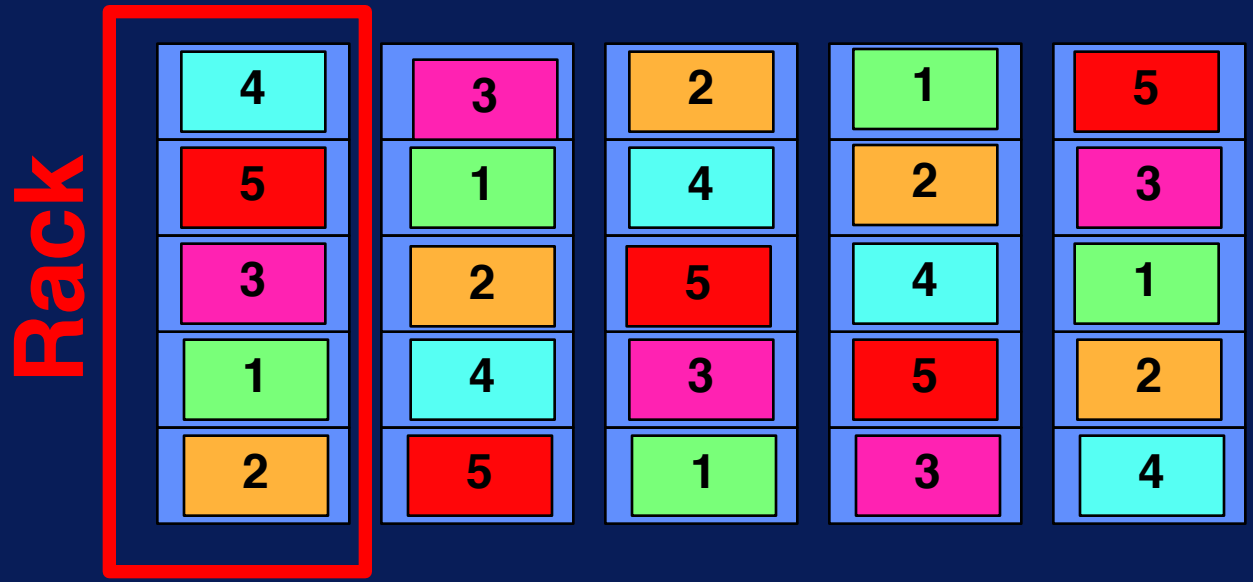

Data

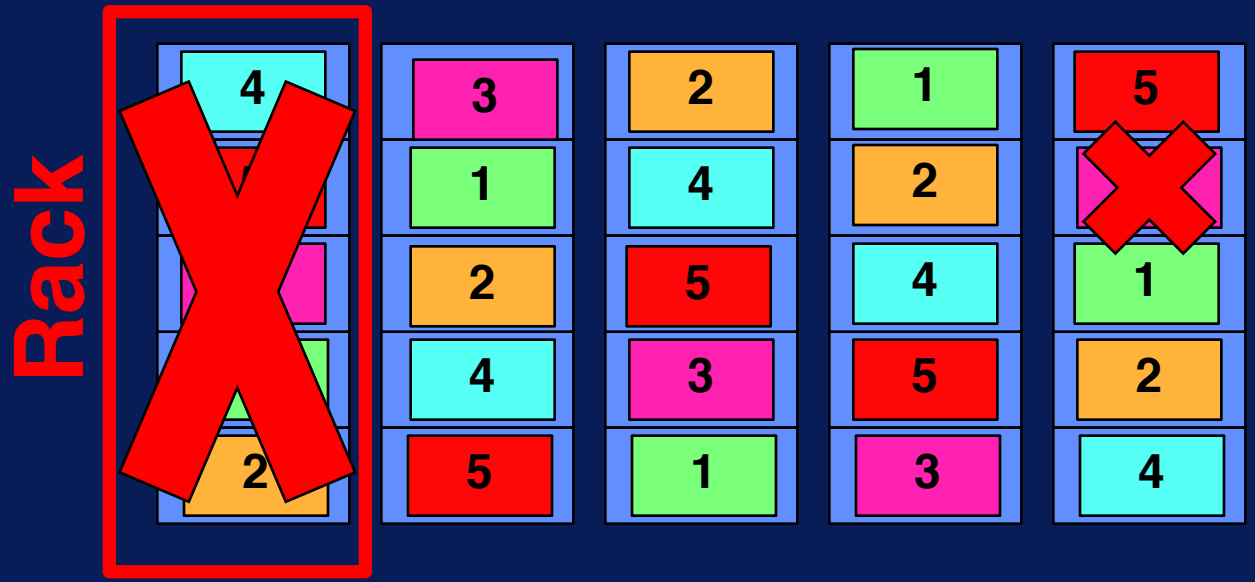

Data

High Concurrency vs. Low Consistency

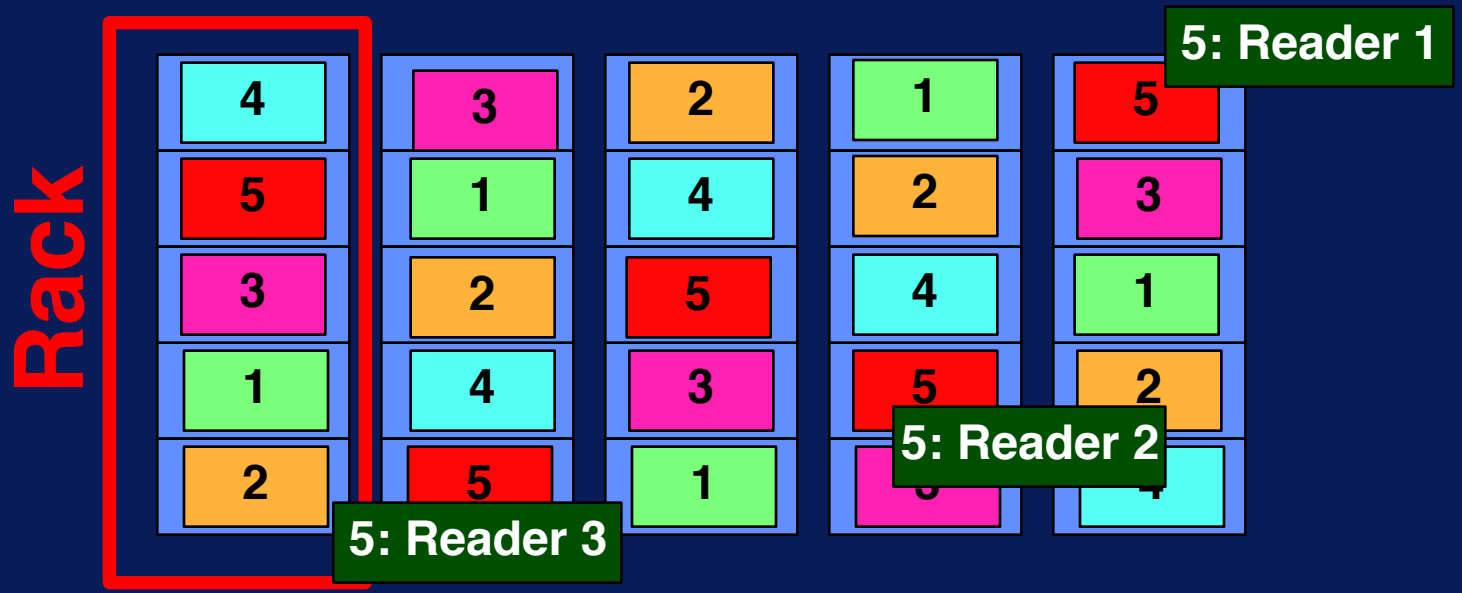

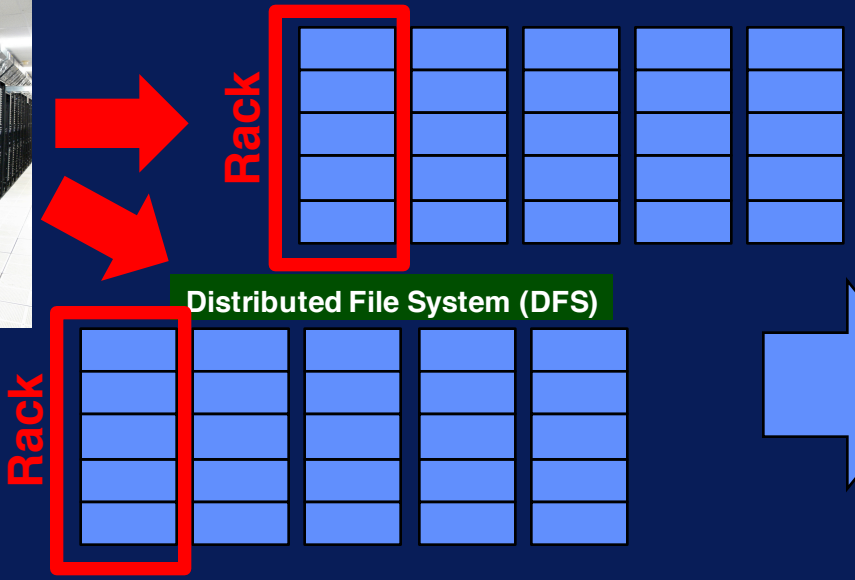

Data partitioning

Data replication

Data scalability

Fault tolerance

High concurrency